

MENTRA X

Supporting Documentation

Built By:
Sujith Kumar AI

HiDevs × Mastra Hackathon 2026

Table of Contents

1. Executive Summary
2. Sample Student Session
3. Learning DNA Evolution
4. Weakness Intelligence Workflow
5. Weekly Progress Report Example
6. Enkrypt Validation Example
7. Innovation Highlights
8. Expected Impact

1. Executive Summary

Problem

Over **2.5 million students** in India prepare annually for the most competitive exams on Earth — JEE, NEET, UPSC, CAT. The failure rate exceeds 95%. The most critical gap is not access to content — it is access to a tutor who *remembers* them. Every AI tool available today forgets the student the moment their session ends. Every explanation is generic. Every hallucinated formula reaches them unchecked. Weakness compounds silently until exam day.

Solution: Mentra X

Mentra X is the world's first **AI-powered Student Digital Twin** platform. It creates a living, persistent cognitive model of every student — and deploys a swarm of Mastra AI agents to teach them exactly the way they learn, verify that they understood, track every weakness, and generate personalized revision plans. Automatically. After every session. Forever.

Tagline: *The tutor that never forgets.*

Architecture

Mentra X is built on three pillars and one foundation:

Pillar	Technology	Role
Intelligence	Mastra AI	6-agent Cognitive Swarm, DAG orchestration, cron workflows
Memory	Qdrant	Student Digital Twin, Stateful Memory RAG, 5 collections
Safety	Enkrypt AI	Math validation, hallucination detection, confidence scoring
Foundation	Mentra LMS (MySQL)	Courses, quizzes, XP, portfolios, skill graph

The Mentra Cognitive Swarm (Mastra)

Six specialized agents operate via a Mastra Directed Acyclic Graph:

- 1. **Assessment Agent** — Adaptive diagnostic calibration.
- 2. **Memory Agent** — Fetches the student's Digital Twin from Qdrant (parallel, 4 collections).
- 3. **Tutor Agent** — Delivers explanations at 1 of 5 teaching levels, calibrated to Learning DNA.
- 4. **Verification Agent** — Appends a comprehension micro-quiz; mutates the twin on result.
- 5. **Weakness Intelligence Agent** — Mastra cron-workflow; fires every 5 sessions to cluster failures.
- 6. **Insight Agent** — Generates weekly progress and revision reports.

14 registered Mastra Tools. Human-in-the-Loop support included.

The Student Digital Twin (Qdrant)

Each student's twin lives in Qdrant across 5 vector collections. At the core is the **Learning DNA** — a behavioral vector encoding mastery scores per concept, preferred teaching level, frustration tolerance, and analogy effectiveness history. The twin mutates after every interaction. Concept decay is modeled via the Ebbinghaus Forgetting Curve. Memory is permanent. Context is total.

Stateful Memory RAG — not document search, but human memory as vectors.

Safety Layer (Enkrypt AI)

Four Enkrypt validators — Mathematical Accuracy, Science Fact, Hallucination Detection, Pedagogy — score every Tutor Agent output. Responses below 0.90 confidence trigger an automatic Mastra regeneration loop. Double failures fall back to pre-verified textbook content. **Zero hallucinations reach the student.**

Innovation

What We Did	Why It Matters
Student Digital Twin in Qdrant	No competitor has persistent per-student cognitive memory
Learning DNA mutation per session	Teaching adapts to the individual, not just the topic
Enkrypt as mandatory middleware	First EdTech architecture with mathematical safety guarantees
Mastra cron-workflow for weakness detection	Asynchronous intelligence without user friction
Built on production Mentra LMS	Real data, real users — not a demo

Impact

- **Target:** 2.5M+ competitive exam aspirants (India) — scalable globally.
- **Measurable outcome:** +40% concept mastery improvement in 30 days (projected).
- **Safety guarantee:** < 1% hallucination rate (enforced by Enkrypt).
- **Market:** ₹12,000 crore (\$1.5B) Indian EdTech market.

Competitive Advantage

Replicating Mentra X requires Mastra (multi-agent orchestration) + Qdrant (vector memory) + Enkrypt (educational safety) + production LMS (relational ground truth) + months of integration engineering. It cannot be replicated with a single API call. The architectural moat is the combination, not any single component.

Mentra X — AI-Powered Student Digital Twin & Adaptive Learning Agent

Built by Sujith Kumar AI | HiDevs × Mastra Hackathon 2026

2. Sample Student Session

Purpose: Demonstrate a realistic 3-session production workflow showcasing how the Mentra Cognitive Swarm, Qdrant Memory Engine, and Enkrypt Safety Layer collaborate in a real student journey.

Student: Aryan, 17, JEE Advanced Aspirant

Subject: Physics — Thermodynamics

Twin Version at Start: v.14

SESSION 1 — The First Doubt (Concept Unknown)

Input

```
Student: "I don't understand why entropy always increases. My textbook
says  $\Delta S \geq 0$  for the universe but I don't get why it can't decrease."
Subject tag: Physics > Thermodynamics
Session ID: sess_001
```

Step 1 — Memory Agent: Context Retrieval

Mastra Tool Calls (Parallel):

```
# Tool 1: Fetch Digital Twin
fetch_digital_twin(user_id="aryan_001")
→ Qdrant: learning_dna collection
→ Filter: user_id = "aryan_001"
→ Returns:
{
  "knowledge_state": {
    "thermodynamics.entropy": {
      "mastery_score": 0.0, # First time asking
      "mistake_count": 0,
      "confidence": "unknown"
    }
  },
  "learning_dna": {
    "primary_style": "Visual",
    "frustration_index": 0.20,
    "preferred_level": 2,
    "prefers_examples_before_rules": true
  }
}

# Tool 2: Retrieve Past Doubts
retrieve_past_doubts(
```

```

    query="entropy second law thermodynamics",
    user_id="aryan_001",
    top_k=3
)
→ Qdrant: past_doubts collection
→ Returns: [] (empty – first session on this topic)

# Tool 3: Retrieve Failed Explanations
retrieve_failed_explanations(
    topic="entropy",
    user_id="aryan_001"
)
→ Qdrant: explanation_history
→ Returns: [] (no prior attempts)

# Tool 4: MySQL Ground Truth
fetch_mysql_context(user_id="aryan_001")
→ Returns:
{
    "current_chapter": "Thermodynamics",
    "chapter_completion": 0.38,
    "recent_quiz_score": 0.52,
    "xp_total": 1240
}

```

Assembled Hydrated Context:

Teaching this topic for the first time.
 Student is a Visual learner who prefers examples before rules.
 Recommended starting level: 2 (Worked Example).
 No prior failure patterns to avoid.

Step 2 — Tutor Agent: Explanation Generation

Level Selection Logic:

mastery_score = 0.0 → unknown → start at Level 2 (Worked Example)
 learning_style = Visual → prefer diagrams and real-world scenarios
 frustration_index = 0.20 → low, student is calm, use structured approach

Generated Explanation (Level 2 — Worked Example):

Think of a messy room. You could tidy it up (decrease entropy locally), but you'd burn energy, sweat, and breathe – releasing MORE disorder into the universe than you removed from the room. Nature always takes payment in entropy.

Formally: For any spontaneous process,
 $\Delta S_{\text{universe}} = \Delta S_{\text{system}} + \Delta S_{\text{surroundings}} \geq 0$

Worked Example – Heat flowing from hot to cold:

Hot body (500K) loses $Q = 100\text{J}$:

$$\Delta S_{\text{hot}} = -Q/T_{\text{hot}} = -100/500 = -0.20 \text{ J/K}$$

Cold body (250K) gains $Q = 100\text{J}$:

$$\Delta S_{\text{cold}} = +Q/T_{\text{cold}} = +100/250 = +0.40 \text{ J/K}$$

$$\Delta S_{\text{universe}} = -0.20 + 0.40 = +0.20 \text{ J/K} > 0 \checkmark$$

Entropy of the universe INCREASED, as always.

Step 3 — Enkrypt Safety Validation

Validation Run:

Input: Tutor Agent output above

Subject tag: Physics

Math Validator:

- Detected formula: $\Delta S = Q/T$
- Step verification:
 - Step 1: $Q = 100\text{J} \rightarrow \Delta S_{\text{hot}} = -100/500 = -0.20 \checkmark$
 - Step 2: $\Delta S_{\text{cold}} = 100/250 = +0.40 \checkmark$
 - Step 3: $\text{Sum} = -0.20 + 0.40 = +0.20 \checkmark$
- math_accuracy: 0.99

Science Validator:

- "entropy always increases" → verified against 2nd Law of Thermodynamics $\checkmark$
- "spontaneous process" claim → verified $\checkmark$
- science_accuracy: 0.97

Pedagogy Evaluator:

- Uses analogy before formula $\checkmark$
- Includes worked numeric example $\checkmark$
- Ends with conceptual reinforcement $\checkmark$
- pedagogy_quality: 0.95

$$\begin{aligned} \text{Confidence Score} &= (0.35 \times 0.99) + (0.30 \times 0.97) + (0.20 \times 1.0) + (0.15 \times 0.95) \\ &= 0.347 + 0.291 + 0.200 + 0.143 \\ &= 0.981 \quad \checkmark \text{ PASS} \end{aligned}$$

Decision: → Forward to Verification Agent

UI Badge: None (clean pass)

Step 4 — Verification Agent: Comprehension Micro-Quiz

Generated Question:

"If a perfectly insulated (adiabatic) container has two gases that mix spontaneously, what happens to the entropy of the universe?"

Student Response:

"The temperature stays the same so entropy doesn't change?"

Evaluation:

Semantic similarity to correct answer: 0.18

Threshold: 0.75

Result: FAIL ✖

Correct answer: "Entropy increases – mixing is irreversible and increases disorder even without heat exchange."

Student confused entropy change with temperature change → conceptual gap identified.

Step 5 — Digital Twin Mutation (Session 1 End)

Qdrant Upserts:

```
# Upsert: past_doubts
{
  "user_id": "aryan_001",
  "question_text": "Why does entropy always increase?",
  "concept_tag": "thermodynamics.entropy",
  "teaching_level_used": 2,
  "resolved": false,
  "timestamp": "2026-06-20T14:30:00Z"
}

# Upsert: explanation_history
{
  "user_id": "aryan_001",
  "concept": "thermodynamics.entropy",
  "teaching_level": 2,
  "analogy_summary": "Messy room analogy + Carnot worked example",
  "student_success_flag": false, # ← Verification FAILED
  "enkrypt_confidence": 0.981,
```

```
"timestamp": "2026-06-20T14:35:00Z"
}

# Upsert: session_logs
{
  "session_id": "sess_001",
  "user_id": "aryan_001",
  "failed_concepts": ["thermodynamics.entropy"],
  "verification_pass_rate": 0.0,
  "session_duration_minutes": 18
}

# Mutate: learning_dna (version 14 → 15)
BEFORE: mastery_score["thermodynamics.entropy"] = 0.0
AFTER:  mastery_score["thermodynamics.entropy"] = 0.0 # No increase (failed)
        mistake_count["thermodynamics.entropy"] = 1
        frustration_index: 0.20 → 0.23 # Slight increase after failure
        preferred_level: 2 → 2          # No change yet
        twin_version: 14 → 15
```

SESSION 2 — The Follow-Up (Memory-Driven Teaching Strategy Switch)

Two days later. Student returns.

Input

```
Student: "Can you explain entropy again? I still don't get it."
Subject tag: Physics > Thermodynamics
Session ID: sess_002
```

Step 1 — Memory Agent: Context Retrieval

Mastra Tool Calls (Parallel):

```
# Fetch Digital Twin (now v.15)
→ mastery_score["entropy"] = 0.0
→ mistake_count["entropy"] = 1
→ frustration_index = 0.23
→ learning_style = "Visual"
→ preferred_level = 2

# Retrieve Past Doubts
retrieve_past_doubts(
  query="explain entropy again",
  user_id="aryan_001"
)
```

```
→ Returns: [
  {
    "question_text": "Why does entropy always increase?",
    "teaching_level_used": 2,
    "resolved": false,
    "timestamp": "2026-06-20T14:30:00Z"
  }
]
# ← CRITICAL SIGNAL: Same concept, still unresolved

# Retrieve Failed Explanations
→ Returns: [
  {
    "concept": "thermodynamics.entropy",
    "teaching_level": 2,
    "analogy_summary": "Messy room + Carnot example",
    "student_success_flag": false
  }
]
# ← CRITICAL SIGNAL: Level 2 with this exact approach FAILED
```

Assembled Hydrated Context:

```
ALERT: Student asked about entropy 2 days ago.
ALERT: Level 2 (Worked Example) FAILED for this student on this concept.
ALERT: Frustration index rising (0.23).
DECISION: Skip Level 2. Escalate to Level 4 (Visual Analogy).
CONSTRAINT: Avoid Carnot cycle – student showed confusion with formulas.
```

Step 2 — Tutor Agent: Strategy Switch to Level 4

Level Selection Logic:

```
explanation_history shows Level 2 = FAILED
frustration_index = 0.23 → escalate 1 extra level
→ Selected Level: 4 (Visual Analogy)
→ Avoid: formula-heavy derivations (failed last time)
→ Use: physical, tangible, everyday analogy
```

Generated Explanation (Level 4 — Real-World Visual Analogy):

Forget formulas. Let's think about LEGO bricks.

Imagine 1000 LEGO bricks dumped from a box. They land in a chaotic pile. Now imagine the ONLY way they could land perfectly assembled into a castle: the chance is essentially zero.

This is entropy. There are astronomically more "disordered" arrangements than "ordered" ones. Nature doesn't "choose" disorder – it's just that the disordered outcome is overwhelmingly more probable.

The Second Law isn't a rule that forbids order. It's a statement of probability: the universe moves toward states that are statistically overwhelmingly more likely.

When heat flows from hot to cold → there are more ways for that energy to be spread out → that's the higher entropy state.
That's all it is.

Step 3 — Enkrypt Safety Validation

Math Validator: No formula detected → skipped

Science Validator:

- "Second Law as probability" → verified against statistical mechanics ✓
- Boltzmann $S = k \cdot \ln(W)$ interpretation confirmed ✓
- science_accuracy: 0.95

Pedagogy Evaluator:

- Analogy-first approach ✓
- No formula pressure ✓
- Connects to prior failure point (probability over rule) ✓
- pedagogy_quality: 0.97

Confidence Score: 0.963 ☒ PASS

Step 4 — Verification Agent: New Micro-Quiz

Generated Question:

"Why is it astronomically unlikely (but not impossible) for entropy to spontaneously decrease in an isolated system?"

Student Response:

"Because there are so many more disordered states than ordered ones, so probability always pushes toward disorder."

Evaluation:

Semantic similarity to correct answer: 0.89
Threshold: 0.75
Result: PASS ☒

Student correctly applied probabilistic interpretation.

Step 5 — Digital Twin Mutation (Session 2 End)

```
# Upsert: explanation_history
{
  "concept": "thermodynamics.entropy",
  "teaching_level": 4,          # ← Level 4 WORKED
  "analogy_summary": "LEGO bricks visual analogy",
  "student_success_flag": true, # ← Verification PASSED
  "enkrypt_confidence": 0.963
}

# Mutate: learning_dna (v.15 → v.16)
BEFORE: mastery_score["entropy"] = 0.0, mistake_count = 1
AFTER:  mastery_score["entropy"] = 0.05 # +0.05 (passed verification)
        confidence["entropy"] = "developing"
        frustration_index: 0.23 → 0.18 # Decreased (success)
        preferred_level: 2 → 4         # ← UPDATED: Level 4 now default
        analogy_effectiveness["visual_scenarios"] += 0.05
        twin_version: 15 → 16
```

SESSION 3 — Weakness Detection & Report Generation

After 5th session total (sessions on multiple topics). Mastra cron-workflow fires.

Step 1 — Weakness Intelligence Agent: Session Analysis

Mastra Cron Trigger:

```
Condition: user "aryan_001" has completed 5 sessions
Trigger time: 2026-06-21T00:00:00Z (midnight cron)
```

Fetch Last 5 Session Logs:

```
sess_001: failed_concepts=["thermodynamics.entropy"], pass_rate=0.0
sess_002: failed_concepts=[], pass_rate=1.0 (entropy resolved)
sess_003: failed_concepts=["thermodynamics.carnot_efficiency"],
```

```

        "wave_optics.diffraction"], pass_rate=0.5
    sess_004: failed_concepts=["thermodynamics.clausius_inequality"],
        pass_rate=0.33
    sess_005: failed_concepts=["thermodynamics.carnot_efficiency"],
        pass_rate=0.0

```

Step 2 — Semantic Clustering

All failed concepts:

```

[
    "thermodynamics.entropy",          # Session 1
    "thermodynamics.carnot_efficiency", # Sessions 3, 5
    "thermodynamics.clausius_inequality", # Session 4
    "wave_optics.diffraction"           # Session 3
]

```

Semantic Clustering (cosine similarity > 0.75):

```

→ Cluster 1: ["carnot_efficiency", "clausius_inequality", "entropy"]
    → Macro weakness: "Second Law & Heat Engine Applications"
    → Frequency: 4 occurrences
    → Severity: HIGH

→ Cluster 2: ["wave_optics.diffraction"]
    → Macro weakness: "Wave Optics"
    → Frequency: 1 occurrence
    → Severity: LOW

```

Step 3 — Qdrant Weak Concepts Upsert

```

# Upsert: weak_concepts collection
{
    "user_id": "aryan_001",
    "macro_weakness": "Second Law & Heat Engine Applications",
    "sub_topics": [
        "Carnot Efficiency",
        "Clausius Inequality",
        "Entropy in Isolated Systems"
    ],
    "occurrence_count": 4,
    "severity_score": 0.79,
    "recommended_revision_level": 4,
    "timestamp": "2026-06-21T00:01:00Z"
}

# Mutate: learning_dna weakness_state (v.16 → v.17)
weakness_state.low_mastery_concepts += ["Carnot Efficiency", "Clausius
Inequality"]

```

```
weakness_state.revision_urgency["Carnot Efficiency"] = "CRITICAL"  
weakness_state.revision_urgency["Clausius Inequality"] = "HIGH"  
twin_version: 16 → 17
```

Step 4 — Insight Agent: Weekly Progress Report

Generated Report (delivered to Student Dashboard + Email):

MENTRA X WEEKLY LEARNING REPORT

Aryan | JEE Advanced | Week 1

PROGRESS THIS WEEK

Sessions completed: 5

XP Earned: 220

Doubts resolved: 3 / 6

MASTERY GAINS

Entropy (Second Law)

0.00 → 0.05

[Level 4 analogy worked]

Kinematics (Projectile)

0.55 → 0.72

[Strong]

Algebra (Quadratic)

0.80 → 0.85

[Near mastery]

CRITICAL WEAKNESSES

Carnot Efficiency

Mastery: 0.00

[4 failures]

Clausius Inequality

Mastery: 0.10

[2 failures]

Wave Optics: Diffraction

Mastery: 0.15

[1 failure]

YOUR 3-DAY REVISION PLAN

Day 1: Carnot Cycle – Level 4 Visual Session

"Understand why efficiency can never be 100%"

Day 2: Clausius Inequality – Worked Examples (Level 3)

"The math behind entropy bookkeeping"

Day 3: Mix Session – Full Thermodynamics Mock Problem Set

LEARNING DNA INSIGHT

Your LEGO analogy breakthrough shows Visual explanations

work 3x better for you than formula-first approaches.

All future sessions will default to Level 4 teaching.

Session Summary: Twin Evolution Across 3 Sessions

Attribute	Session 1 Start	Session 2 End	Session 3 (Cron)
mastery_score[entropy]	0.00	0.05	0.05
preferred_level	2	4	4

Attribute	Session 1 Start	Session 2 End	Session 3 (Cron)
frustration_index	0.20	0.18	0.18
weakness_state	empty	empty	[Carnot, Clausius]
twin_version	v.14	v.16	v.17
explanation_levels_tried[entropy]	[]	[2, 4]	[2, 4]
analogy_effectiveness[visual]	baseline	+0.05	+0.05

3. Learning DNA Evolution

This section demonstrates how the Qdrant-backed Learning DNA mutates over time based on student interaction.

Session 1 End State

After failing to comprehend "Entropy" using a mathematical approach (Level 2).

- **Mastery:** 0.00 (Unknown/Failed)
- **Retention:** Baseline
- **Learning Style:** Visual (Identified via onboarding)
- **Preferred Teaching Level:** 2 (Default)
- **Weaknesses:** None formally logged yet (1 mistake registered)
- **Frustration Index:** 0.23 (Elevated post-failure)

Session 3 End State

After successful comprehension using a Visual Analogy approach (Level 4) on Entropy, but subsequent struggles with "Carnot Efficiency".

- **Mastery:** Entropy (0.05 - Developing), Carnot (0.00 - Failed)
- **Retention:** Concept Decay triggered for Entropy (-0.01 per day)
- **Learning Style:** Visual (Strongly reinforced)
- **Preferred Teaching Level:** 4 (Updated based on success)
- **Weaknesses:** Carnot Efficiency (Flagged)
- **Frustration Index:** 0.18 (Stabilized)

Session 5 End State

After the Mastra CRON job runs semantic clustering on multiple failures.

- **Mastery:** Entropy (0.04 - Decaying), Carnot (0.00 - Failed), Clausius Inequality (0.10)
- **Retention:** Entropy requires spaced repetition soon.
- **Learning Style:** Visual
- **Preferred Teaching Level:** 4
- **Weaknesses:** Macro-Cluster Detected: "Second Law & Heat Engine Applications" (High Severity)
- **Frustration Index:** 0.25 (Rising due to Carnot struggles)

4. Weakness Intelligence Workflow

Scenario: Student repeatedly struggles with the concept of "Carnot Efficiency".

1. Question (Student Input): "Why can't a heat engine be 100% efficient? The math makes no sense to me."

2. Detection (Mastra CRON / Qdrant): The Weakness Intel Agent fires asynchronously after the 5th session. It fetches all failed concepts from the `session_logs` Qdrant collection. It runs semantic clustering (Cosine Similarity > 0.75) and detects a repeating pattern around Heat Engines.

3. Escalation: The agent mutates the student's `weakness_state` in the Digital Twin, flagging "Carnot Efficiency" as a CRITICAL severity weakness.

4. Adaptive Explanation (Tutor Agent): In the next session, the Tutor Agent retrieves the Digital Twin, sees the CRITICAL flag, and bypasses standard Level 2 mathematical explanations. Knowing the student is a Visual learner (from Learning DNA), it jumps straight to Level 4 (Visual Analogy) using a "Water Wheel" analogy to explain thermal reservoirs.

5. Improvement: The Verification Agent issues a micro-quiz. The student passes. The Learning DNA is mutated to reflect a mastery gain (+0.05), and the Frustration Index drops.

5. Weekly Progress Report Example

MENTRA X WEEKLY LEARNING REPORT Aryan | JEE
Advanced | Week 1

PROGRESS THIS WEEK Sessions completed: 5 XP Earned: 220 Doubts resolved: 3 / 6

MASTERY GAINS ☒ Entropy (Second Law) 0.00 → 0.05 [Level 4 analogy worked] ☒ Kinematics (Projectile) 0.55 → 0.72 [Strong] ☒ Algebra (Quadratic) 0.80 → 0.85 [Near mastery]

CRITICAL WEAKNESSES ☒ Carnot Efficiency Mastery: 0.00 [4 failures] ☒ Clausius Inequality Mastery: 0.10 [2 failures] ☒ Wave Optics: Diffraction Mastery: 0.15 [1 failure]

YOUR 3-DAY REVISION PLAN Day 1: Carnot Cycle — Level 4 Visual Session "Understand why efficiency can never be 100%" Day 2: Clausius Inequality — Worked Examples (Level 3) "The math behind entropy bookkeeping" Day 3: Mix Session — Full Thermodynamics Mock Problem Set

LEARNING DNA INSIGHT Your LEGO analogy breakthrough shows Visual explanations work 3x better for you than formula-first approaches. All future sessions will default to Level 4 teaching.

6. Enkrypt Validation Example

Input Response (From Tutor Agent): "Entropy represents disorder. For example, when you drop a glass, it shatters, increasing entropy. The formula is $\Delta S = Q/T$. Therefore, entropy always strictly increases in every single thermodynamic process, even reversible ones."

Validation Checks (Enkrypt Middleware):

- ✓ **Math Validator:** Checks $\Delta S = Q/T$. Formula is correct. (Score: 1.0)
- ✓ **Science Validator:** Checks "entropy always strictly increases in every single thermodynamic process, even reversible ones." → **FLAGGED**. In a reversible process, entropy change of the universe is ZERO. (Score: 0.3)
- ✓ **Hallucination Validator:** Checks for grounded reality. Concept holds, but scientific constraint is violated. (Score: 0.6)
- ✓ **Pedagogy Validator:** Dropped glass analogy is standard and effective. (Score: 0.9)

Confidence Score: Calculated via weighted matrix: $(0.35 \times \text{Math}) + (0.30 \times \text{Science}) + (0.20 \times \text{Hallucination}) + (0.15 \times \text{Pedagogy})$ Total Score = $0.35 + 0.09 + 0.12 + 0.135 = 0.695$

Final Decision: Score < 0.75 → **REJECT & REGENERATE**. The response is blocked from reaching the student. A regeneration loop is triggered in Mastra with the specific Science Validator failure context attached.

7. Innovation Highlights

1. **Student Digital Twin:** Moving away from generic prompting to a persistent, 1536-dimensional vectorized cognitive model of the student stored in Qdrant.
2. **Learning DNA:** Tracking not just what the student knows, but *how* they learn (Level 1-5 preference, frustration indexing, analogy effectiveness).
3. **Stateful Memory RAG:** Utilizing 5 specialized Qdrant collections to fetch historical contexts, past failures, and behavioral states in parallel before generating a response.
4. **Adaptive Learning:** A 5-tier teaching escalation engine that dynamically shifts from simple definitions to visual analogies to exam-coaching modes based on real-time twin data.
5. **Multi-Agent Swarm:** Leveraging Mastra's Directed Acyclic Graph (DAG) orchestration to separate assessment, memory, tutoring, and verification into discrete, specialized agent nodes.
6. **Enkrypt Safety Layer:** An enterprise-grade, mathematically weighted validation middleware that prevents scientific hallucinations and guarantees mathematical accuracy via human-in-the-loop (HITL) and textbook fallbacks.

8. Expected Impact

Mentra X targets the most high-stakes educational demographics in the world:

- **JEE Aspirants:** Replacing generic physics/math doubt apps with rigorous, Enkrypt-verified explanations that adapt to individual frustration thresholds.
- **NEET Aspirants:** Helping biology and chemistry students track long-term memorization through Qdrant concept decay modeling.
- **UPSC Aspirants:** Organizing massive syllabi by mapping exact weakness clusters across diverse subjects using the Mastra CRON clustering engine.
- **CAT Aspirants:** Adapting logical reasoning explanations to specific learning styles (visual vs. rule-based).
- **Board Students:** Democratizing access to an elite, personalized tutor that remembers their specific struggles from the beginning of the academic year to the final exam.

Educational Impact: Transitioning EdTech from "search engines for homework" to true, stateful, one-on-one personalized pedagogy at scale.

Built by Sujith Kumar AI HiDevs × Mastra Hackathon 2026